

EMU 1.0

Electro-Mechanical Unit for industrial applications

Engineer's Manual

Revision A



HTSP

Department of Computing
and Analytical Engines

EMU 1.0 Engineer's manual

Revision A

Tahlia

Table of Contents

Overview of computer capabilities	5
Functional diagram	5
Component description, interface	6
Tape reader	6
Memory	6
Mathematical unit	6
Device I/O interface.....	6
Serial port device	7
Serial port character set	7
Clock device	7
Instruction encoding	8
Arithmetic and logic	9
Comparison	9
Shift/rotate by immediate	10
Indirect memory access	10
Fixed-point multiply	10
Control flow	11
Subroutine example	11
Device I/O	12

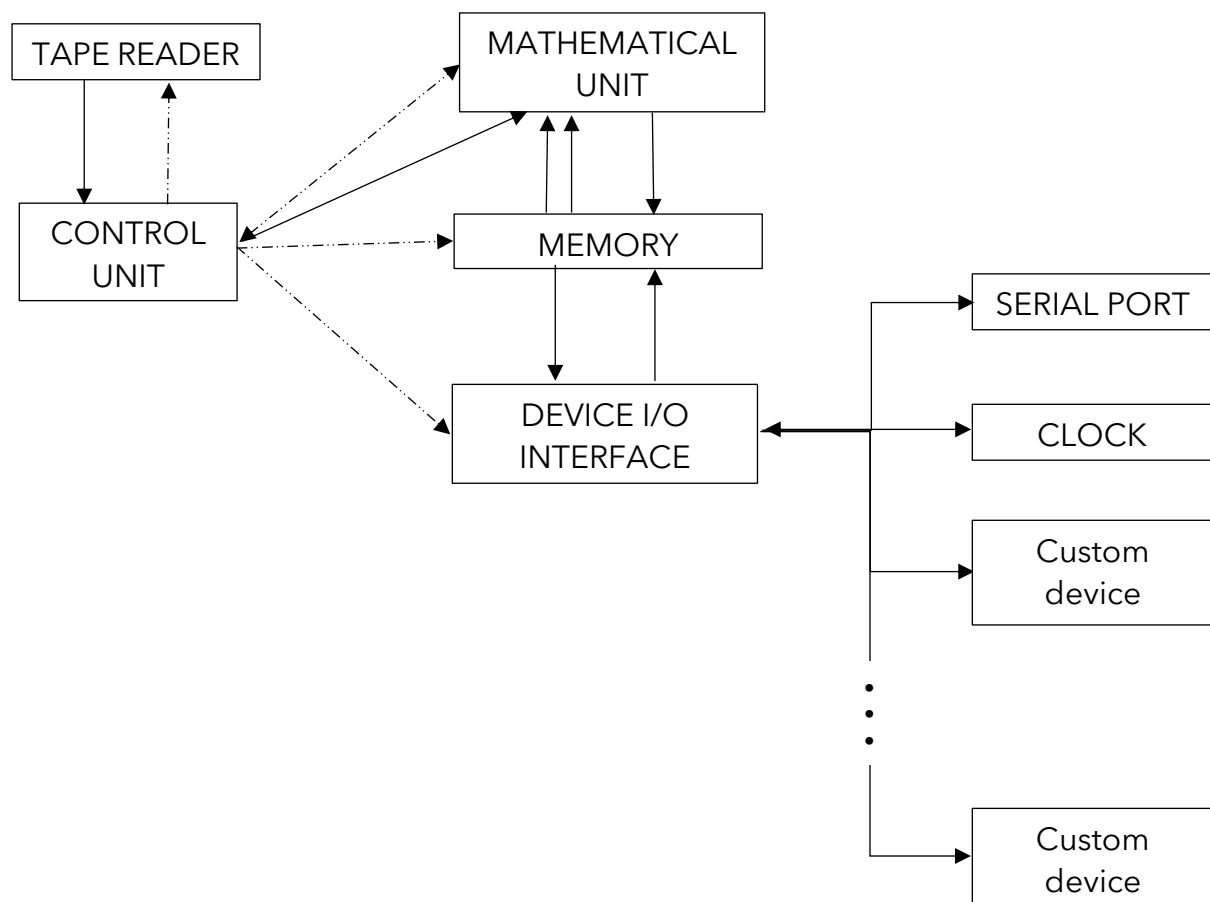
Overview of computer capabilities

The EMU 1.0 is a Harvard architecture 6-bit relay computer. It has 63 6-bit words of working memory, addressable either directly, as general-purpose registers, or indirectly, as Random Access Memory. Its program memory is stored on 24 row punched tape and is unlimited in size. EMU 1.0 can perform a variety of arithmetic and bitwise operations, including fixed point multiplication, and it has an extensible device I/O subsystem.

- Harvard architecture
- 6-bit arithmetic, logic, fixed point multiplication
- 63 words of working memory
- Unlimited program memory; punched tape, 24 rows
- Up to 60 custom I/O devices

Functional diagram

Dashed lines represent control signals. Solid lines represent data paths.



Component description, interface

Tape reader

The tape reader outputs the instruction currently under the read head to the control unit. The control unit can send "move forward" and "move backward" signals to the tape reader, which will actuate the tape accordingly by one step.

Memory

The memory bank has 64 addressable slots, selected via control lines. It interfaces with the other components via 3 data output latches: 2 for the math unit inputs, one for outputting data to devices, and 2 data inputs for storing math results and receiving data from devices.

In EMU 1.0 assembly, memory slots are referred to using the r prefix: r0, r1, ..., r63.

Slot 0 will always equate to the value 0 and writing to it has no effect. Slots 1-63 are general purpose and can be used in any instruction.

Control signals: address select (6 bits), output operand 1, output operand 2, output io, input math, input io.

Mathematical unit

The mathematical unit encompasses multiple circuits for the various operations:

- Arithmetic and logic circuit (add, sub, cmp, xor, or, and)
- Shift, rotate circuit (shi, roi, shl, shr, mul)
- Multiplier circuit (mul)

The left operand is always received from the memory component, but the right operand may be either data from the memory or an immediate value received from the instruction decoder. The shifter may either operate on the value of the first operand, or on the output of the multiplier circuit. The ALU always outputs condition flags to the control unit, which are used to set the condition flag during the comparison instructions.

Control signals: immediate b, output select (0 = ALU, 1 = shifter/multiplier); (To ALU) subtract, operation select (0 = add/sub/cmp, 1 = xor, 2 = or, 3 = and); (To shifter) rotate, direction, multiply; (To multiplier) signed.


Output signals: Overflow, Zero, Carry, Sign.

Device I/O interface

On each I/O operation a word is sent to the selected device and a word is received and stored to memory.

Control signals: device select (6 bits), activate.

★ Many tapes are looped !!



Serial port device

The serial port facilitates communication with other computers and exposes 3 device addresses:

Device 0: SERIAL_INCOMING - Reads the number of buffered incoming words, clamped to max 63.

Device 1: SERIAL_READ - Reads the next buffered word, -1 if there are none.

Device 2: SERIAL_WRITE - Writes send a word to the remote machine.

Serial port character set

Communications over the serial port conventionally use the following charset.

	00	01	02	03	04	05	06	07	10	11	12	13	14	15	16	17
00+	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
20+	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V
40+	W	X	Y	Z	SP	+	-	*	/	<	=	>	()	[]
60+	{	}	#	\$	_	?		^	&	!	~	,	.	:	LF	

The last one is an invalid character, **SP** is the space character, **LF** is a line feed.

Clock device

Timing is a strict necessity in control applications, as well as user interfaces. The clock device exposes two device addresses with the following functionality:

Device 3: CLOCK_LO_CS - Reads least significant 6 bits of centisecond counter.

Device 4: CLOCK_HI_CS - Reads most significant 6 bits of centisecond counter.

Writing a nonzero value to either device resets the counter to 0. The 12-bit counter does not overflow and wrap around, instead it is clamped at a max value of 4095 = 40.95 seconds.

Upon computer startup, the value of the counter is automatically set to 0.

?A: seems like octal?!

Digitisation project?
24 bit inst → 4 base 64
symbols?

Instruction encoding

The 24 instruction bits are split into 4 6-bit parts, from now on named **OPC, A, B, C**.


OPC = 0 is an invalid instruction and in turn halts the computer. The other 63 OPC values encode one of the 21 instruction codes (denoted **op**) and a ternary condition code: $OPC = 21 * \text{condition} + \text{op} + 1$. The condition code influences execution as follows:

condition	Assembly text representation	Description
0	(none)	Execute unconditionally
1	+	Execute if the condition flag is true
2	-	Execute if the condition flag is false

For instance, the following EMU 1.0 assembly and C pseudocode are equivalent:

cmpeq r2, r3	flag = r2 == r3;
+ add r4, r5, r6	if(flag) r4 = r5 + r6;
add r5, r5, 011	r5 = r5 + 9;
- sub r4, r5, r6	if(!flag) r4 = r5 - r6;

op = 20 is reserved for future extensions. Executing it should halt the computer.
op 0 through 19 are defined in the following pages.



Arithmetic and logic

op	A	B	C	Assembly	Description
000	rd	ra	rb	add rd, ra, rb	Add ra and rb into rd .
001	rd	ra	ib	add rd, ra, ib	Add ra and immediate ib into rd .
002	rd	ra	rb	sub rd, ra, rb	Subtract rb from ra and store into rd .
004	rd	ra	rb	or rd, ra, rb	Bitwise or ra and rb into rd .
005	rd	ra	ib	or rd, ra, ib	Bitwise or ra and immediate ib into rd .
006	rd	ra	rb	xor rd, ra, rb	Bitwise xor ra and rb into rd .
007	rd	ra	ib	xor rd, ra, ib	Bitwise xor ra and immediate ib into rd .
010	rd	ra	rb	and rd, ra, rb	Bitwise and ra and rb into rd .
011	rd	ra	ib	and rd, ra, ib	Bitwise and ra and immediate ib into rd .
013	rd	ra	rb	shl rd, ra, rb	Logical* shift ra left rb bits into rd .
014	rd	ra	rb	shr rd, ra, rb	Logical* shift ra right rb bits into rd .

* Logical shifts fill the undefined bits with zeros. There is no arithmetic right shift.

Arithmetic and logic instructions only change **rd**, they do not alter any other part of the computer state.

Comparison

op	A	B	C	Assembly	Description
003	000+cc	ra	rb	cmpCM ra, rb	Compare ra and rb .
003	020+cc	ra	ib	cmpCM ra, ib	Compare ra and immediate ib .
003	030+cc	ia	rb	cmpCM ia, rb	Compare immediate ia and rb .

Comparison instructions do not change the contents of the memory, but they set the Control Unit's conditional flag according to their result.

The comparison mnemonics **CM** and their numerical values **cc** are as follows:

cc	CM	Description	Logical expression
0	tr	True	1
1	fa	False	0
2	eq	Equal	ZF = 1
3	ne	Not equal	ZF = 0
4	sl	Signed less than	SF xor OF = 1
5	sg	Signed greater than	SF xor OF = 0 and ZF = 0
6	ul	Unsigned less than	CF = 1
7	ug	Unsigned greater than	CF = 0 and ZF = 0

"Less/greater than or equal" comparisons can be implemented by inverting the condition and swapping the comparison inputs. (e.g. "**a** <= **b**" equivalent to "**b** > **a**")

Octal again

★ Undocumented?!?
003 010+cc rb ra

why?!
crkey

Shift/rotate by immediate

op	A	B	C	Assembly	Description
012	rd	ra	000+ib	shl rd, ra, ib	Logical shift ra left immediate ib bits into rd .
012	rd	ra	010+ib	shr rd, ra, ib	Logical shift ra right immediate ib bits into rd .
012	rd	ra	020+ib	sar rd, ra, ib	Arithmetic shift ra right immediate ib bits into rd .
012	rd	ra	030+ib	rol rd, ra, ib	Rotate ra left immediate ib bits into rd .

Limits: $0 \leq \text{ib} < 8$.

Right rotations can be achieved by their complementary left rotation. Just like the other arithmetic instructions, these only change the destination memory slot.

Indirect memory access

C: VFIJ
EMU: CV+IJ

op	A	B	C	Assembly	Description
015	rd	ra	ib	ld rd, [ra+ib]	Load value at address (slot ra + immediate ib) mod 64 into rd .
016	rs	ra	ib	st [ra+ib], rs	Store rs at address (slot ra + immediate ib) mod 64.

EMU1.0 assemblers should interpret **[ra]** as **[ra+0]** and **[ib]** as **[r0+ib]**.

Fixed-point multiply



op	A	B	C	Assembly	Description
023	rd	ra	000+pr	fmu/pr rd, ra	Multiply unsigned fixed-point rd and ra into rd , with pr fractional bits.
023	rd	ra	020+pr	fms/pr rd, ra	Multiply signed fixed-point rd and ra into rd , with pr fractional bits.

Limits: $0 \leq \text{pr} < 16$.

Multiplies **rd** and **ra** to an internal 12-bit value, shifts right (logical for unsigned, arithmetic for signed) by **pr** bits, storing the final 6 least significant bits to **rd**.

Using **pr** = 6, one can also obtain the upper half of a multiplication result, disregarding its interpretation as a fixed-point value.

Fun: implement some
3D wireframe rendering

Control flow

op	A	B	C	Assembly	Description
017	1a	1b	1c	lbl lab, 1c	Declare jump label lab , lc .
020	1a	1b	rc	jup lab, rc	Jump upwards to label lab , rc .
021	1a	1b	rc	jdn lab, rc	Jump downwards to label lab , rc .

Limits: $0 \leq \text{lab} = 64 \cdot 1a + 1b < 4096$

EMU1.0 assemblers should allow the **lc** or **rc** values to be omitted, case in which they default to **0** and **r0**, respectively.

EMU 1.0 uses a system of designated jump targets to be able to branch without a program counter. Each jump target has an 18-bit identifier given by A, B, C, grouped into **lab** and **lc**. When jumping, the Control Unit scrolls the tape in the hinted direction until it hits a corresponding **lbl** instruction. That is, one which matches both the jump **lab**, as well as **lc** with the value of memory slot **rc**.

For normal jumps, **lc** and **rc** should be zero. Subroutine returns may be implemented by jumping to a return label and a designated "return pointer"-esque memory slot.

Label instructions may be conditional, case in which they will be ignored by the Control Unit when searching for a jump target if the condition dictates so.

Subroutine example

A subroutine at label **01000** is called from multiple places. It returns to the labels **01001**. To differentiate them, each call site uses a different **lc**, which is passed to the subroutine via the **r63** slot.

First call site

```
add r63, r0, 1 # Set "return identifier" to 1
jdn 01000      # Jump to subroutine
lbl 01001, 1   # Return location with identifier 1
```

Second call site

```
add r63, r0, 2 # Set "return identifier" to 2
jdn 01000      # Jump to subroutine
lbl 01001, 2   # Return location with identifier 2
```

...

```
lbl 01000      # Subroutine entry label
# ...          # Omitted subroutine body; does not trash r63
jup 01001, r63 # Jump to whichever return label matches r63
```

Hah!

Will always find lbl when the tape is looped

Device I/O

op	A	B	C	Assembly	Description
022	rd	ix	rs	io rd, ix, rs	Send rs to device ix and receive rd .

EMU 1.0 assemblers should interpret allow either **rd** or **rs** to be omitted, interpreting partial **io** instructions as follows:

```
io rd, ix      =   io rd, ix, r0
io ix, rs      =   io r0, ix, rs
io ix          =   io r0, rx, r0
```

The Device I/O interface enables word-by-word communication with the 64 devices, selected by immediate **ix**. Each **io** invocation sends the value of **rs**, activates the device, gets a response value, and stores it to **rd**. The default **Serial Port** and **Clock** devices and their protocols are described in their respective sections. Other special-purpose devices may be installed on the 59 remaining device slots.

Accessing a device that does not exist/is not connected halts the computer.

Digitise todo:

✓ Mandel Flag

- Robot arm controller
- ELIGMA decypher
- Business manager database
- Blotto Box impl.